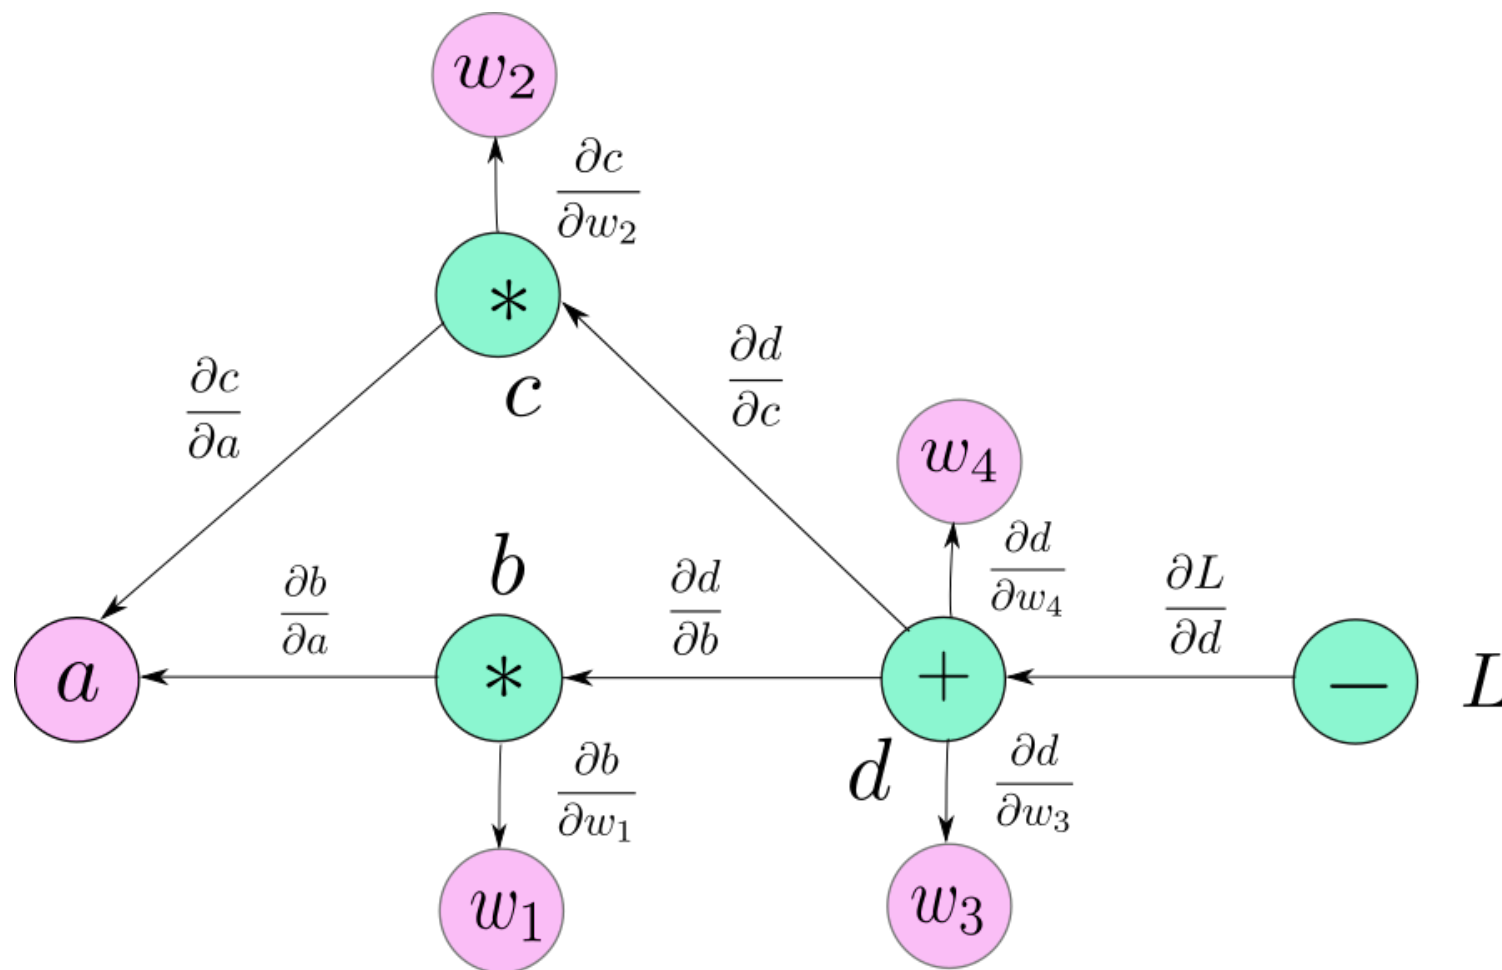


PyTorch Primer - v2

Fundamentals of GenAI
Autograd and Sequential



PyTorch Autograd (autodiff) In a Nutshell



$$\frac{\partial L}{\partial w_4} = \frac{\partial L}{\partial d} * \frac{\partial d}{\partial w_4}$$

$$\frac{\partial L}{\partial w_3} = \frac{\partial L}{\partial d} * \frac{\partial d}{\partial w_3}$$

$$\frac{\partial L}{\partial w_2} = \frac{\partial L}{\partial d} * \frac{\partial d}{\partial c} * \frac{\partial c}{\partial w_2}$$

$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial d} * \frac{\partial d}{\partial b} * \frac{\partial b}{\partial w_1}$$

PyTorch tensor attributes support gradient calculations

```
x = pt.tensor(...)
```

```
x.data
```

```
x.is_leaf = True
```

```
x.requires_grad =  
False
```

```
x.grad = None
```

```
x.grad_fn = None
```

Tensor value(s), such
that `x == x.data`

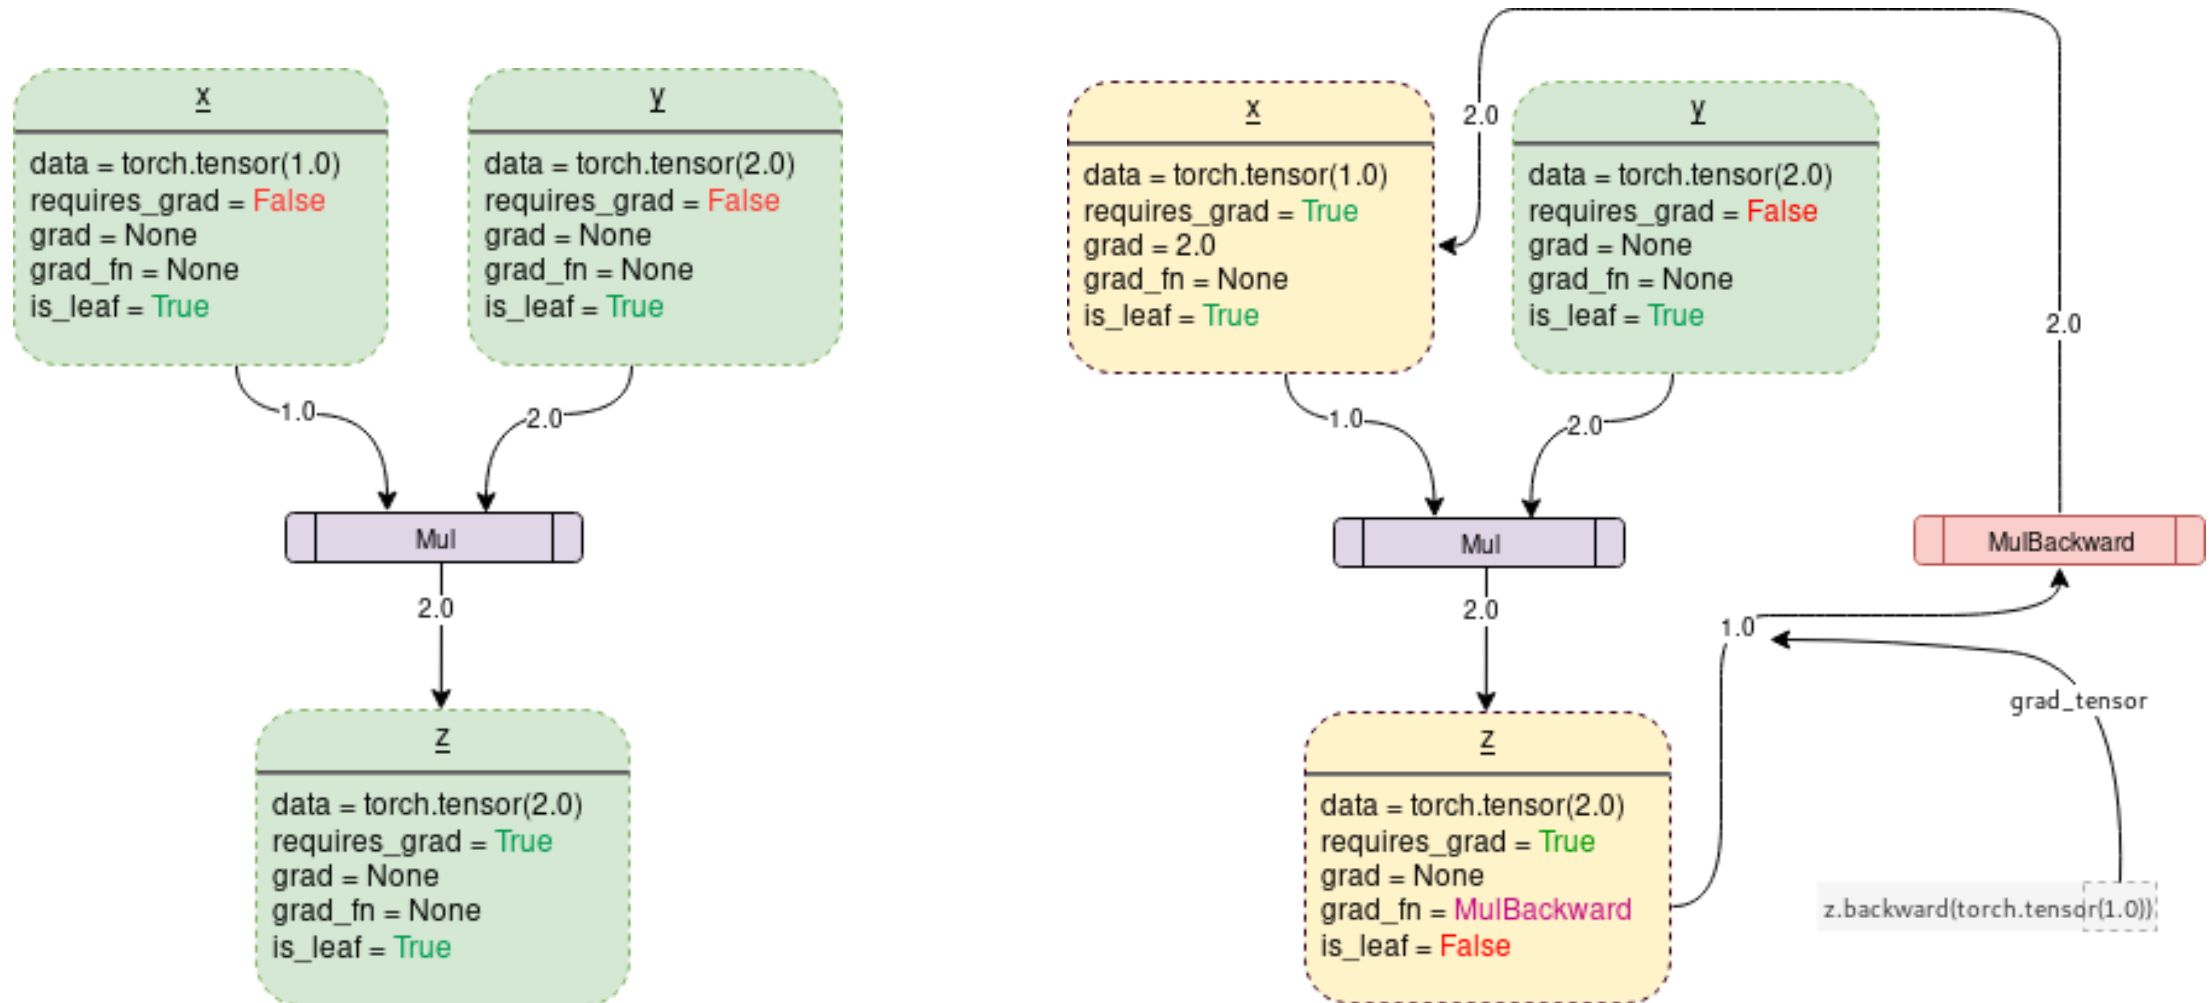
True unless created from a
tensor with `requires_grad=True`

False when `is_leaf = True`

Gradient value if exists

Function that knows how
to perform `backward()`


PyTorch disables tensor autodiff (autograd) by default



Demo

PyTorch AutoGrad


- ▶ Experiment with autograd



Demo

PyTorch AutoGrad

- ▶ Experiment with autograd





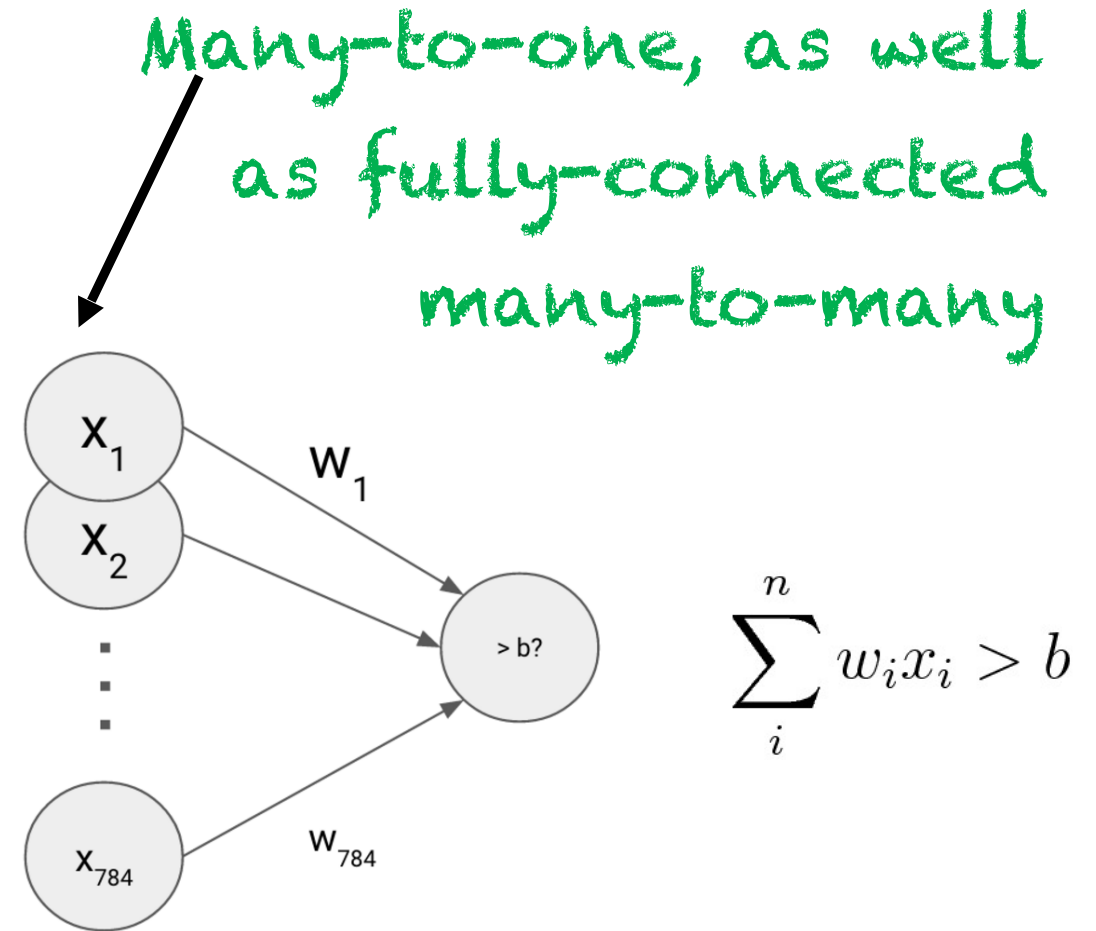
Layers



 **ata Trainers**

- **Available Layers:**

- **Linear**
- Convolution
- Padding
- Pooling
- Normalization
- Dropout
- Recurrent
- Embedding



`nn.Linear()` is a fully connected layer, every input to every output

```
from torch import nn
model = nn.Linear(#_inputs,
#_outputs,
bias = True)

model.parameters()
model.weight, model.bias

model.zero_grad()

model.weight.grad, model.bias.grad
```

Bias by default

Weights
& bias

Zero out
gradients

Access

gradient
values


`model(X)` ← `forward()`



Exercise

Binary Classifier with nn.Linear

- Use the nn package



To read sharded CSV files, use Pandas **concat** and then decode the CSV into features and targets for the training **TensorDataset**

```
from torch.utils.data import TensorDataset

df = pd.concat(
    pd.read_csv(file) for file in Path('data/').glob('part-*.csv')
)
...

X = pt.tensor(train_df[FEATURES].values)
y = pt.tensor(train_df[TARGET].values)

train_ds = TensorDataset(y, X)
```

TensorDataset is used to instantiate an enumerable DataLoader, which supports data batching and indexing

```
from torch.utils.data import DataLoader

train_ds = TensorDataset(y, x)

train_dl = DataLoader(train_ds, batch_size=BATCH_SIZE)


for epoch in range(EPOCHS):
    for batch_idx, batch in enumerate(train_dl):
        y, x = batch
```



Exercise


Dataset and DataLoaders

- Use the data package





torch.nn.Sequential helps organize neural nets with many layers

```
model = nn.Sequential(  
    nn.Linear(5, 5),  
    nn.ReLU(),  
    nn.Linear(5, 5),  
    nn.ReLU(),  
    nn.Linear(5, 5)  
)  
model(X)
```




Softmax helps deal with non-binary (multivariate) target values

input vector
(pixel data)



output vector
(probability
estimate)

$$p(y = j | \mathbf{x}) = \frac{\exp(\mathbf{w}_j^T \mathbf{x} + b_j)}{\sum_{k \in K} \exp(\mathbf{w}_k^T \mathbf{x} + b_k)}$$



Negative log likelihood is the loss for **softmax** outputs

$$LogLoss = \sum_{(\mathbf{x}, y) \in D} -y \log(y') - (1 - y) \log(1 - y')$$

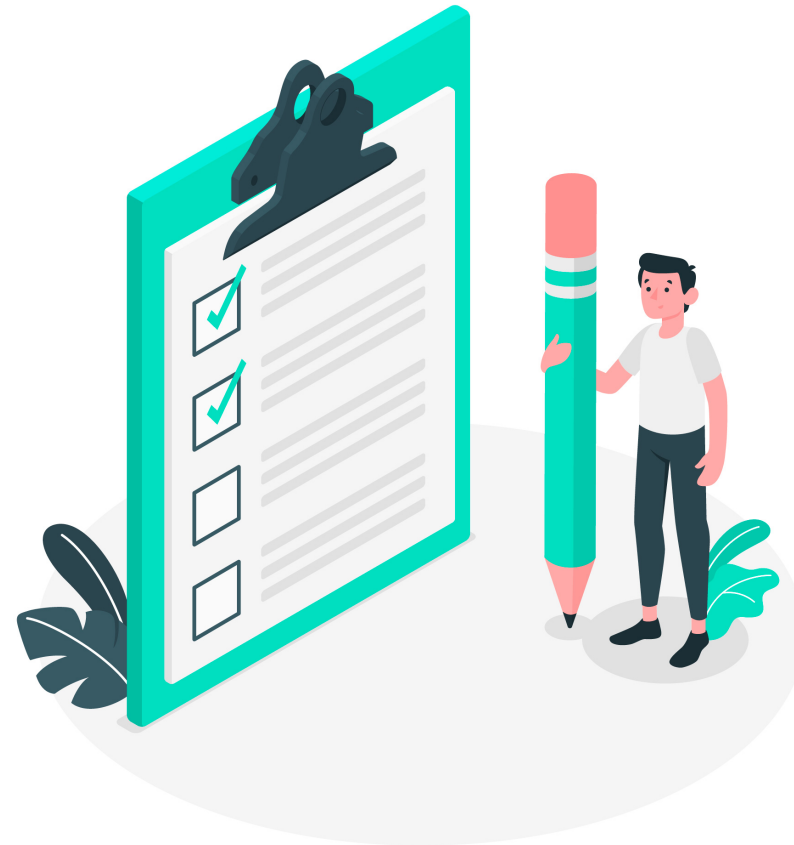
`torch.nn.functional.log_softmax() + nll_loss() = cross_entropy()`



Demo

Sequential Networks

► Use nn.Sequential



Lab

Sequential Networks

► Use nn.Sequential

